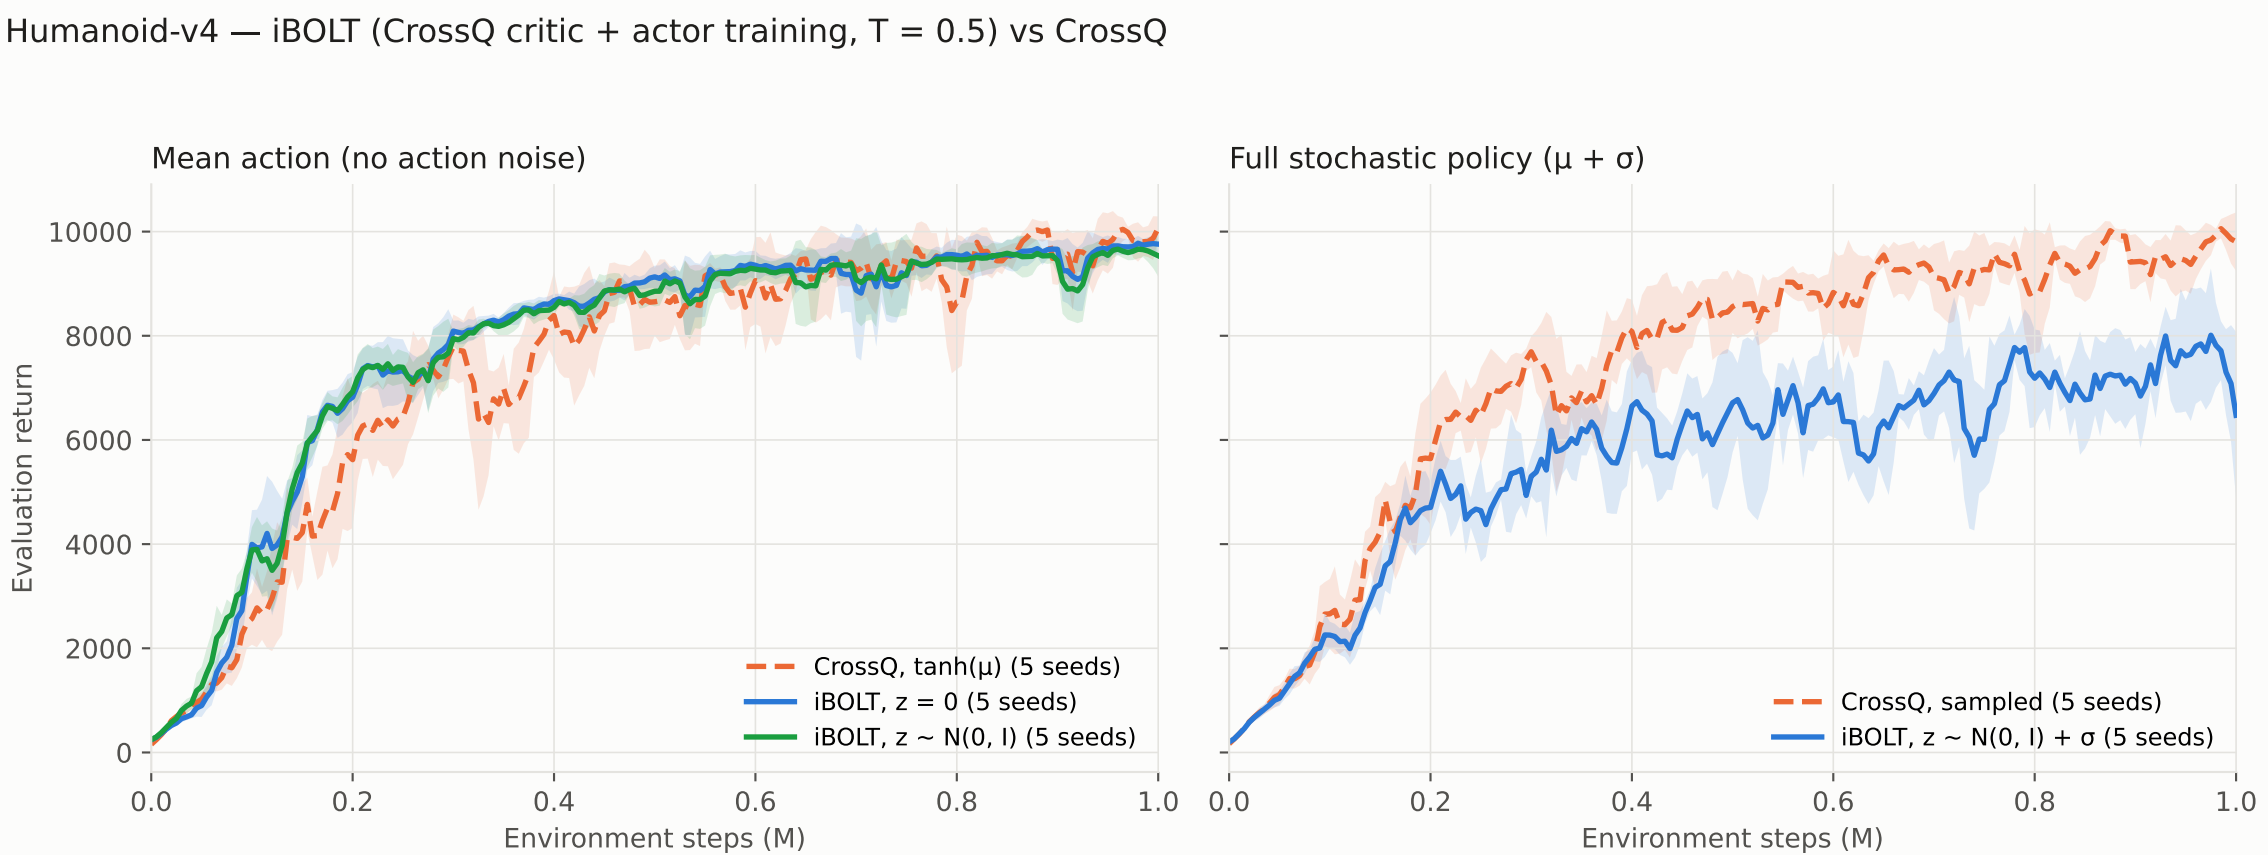
Humanoid-v4 — iBOLT (CrossQ critic + actor training, T = 0.5) vs CrossQ
Mean action (no action noise)
Full stochastic policy (μ + σ)
Evaluation return
0
2000
4000
6000
8000
10000
0.0
0.2
0.4
0.6
0.8
1.0
Environment steps (M)
CrossQ, tanh(μ) (5 seeds)
iBOLT, z = 0 (5 seeds)
iBOLT, z ~ N(0, I) (5 seeds)
CrossQ, sampled (5 seeds)
iBOLT, z ~ N(0, I) + σ (5 seeds)